



BUILDING A REAL-TIME STOCK MARKET DATA PIPELINE WITH KAFKA & AWS

By Winardi





HI, I AM WINARDI

Passionate about building reliable data pipelines and scalable data systems, with a strong foundation in statistics. Focused on data engineering to ensure clean and trustworthy data for better decision-making.



Course License :

- Data Engineer –[DSAREA \(Oct 2025\)](#)
 - Learn AI Basics @Dicoding
 - Getting Started Programming in Python @Dicoding
 - Learn Basic Structured Query Language (SQL) @Dicoding
 - Excel Intermediate @Koding Studio
 - Data Science Bootcamp @Kelas Wok by Kelas.com
 - Microsoft Power BI @MySkill
- 

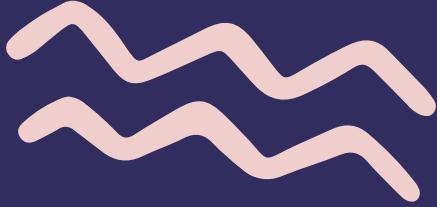
Contact Info :

Email : awinardi1004@gmail.com

LinkedIn : [winardi-](#)

GitHub : [awinardi1004](#)

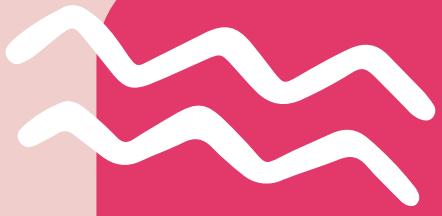
BACKGROUND



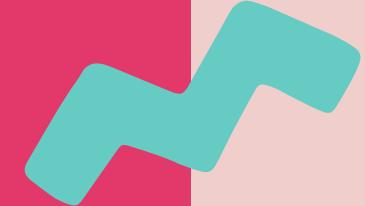
Stock prices move every second, and processing this fast-changing data requires a system that is both reliable and real-time. Many existing solutions still rely on batch processing, which often delays insight generation.

This project was developed to build a real-time data pipeline capable of collecting, streaming, storing, and analyzing stock market data using Kafka and AWS services.





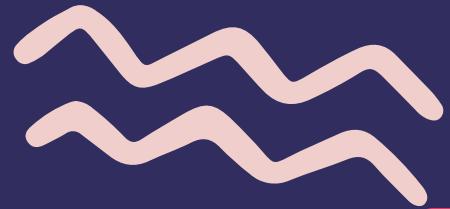
PROBLEM STATEMENT



How can we design a system that continuously handles streaming stock data end to end—from API ingestion, Kafka streaming, storage in S3, to fast querying in Athena?

The main challenges include the speed of incoming data, dynamic data structures, and the need for efficient storage that remains easy to query.

OBJECTIVES



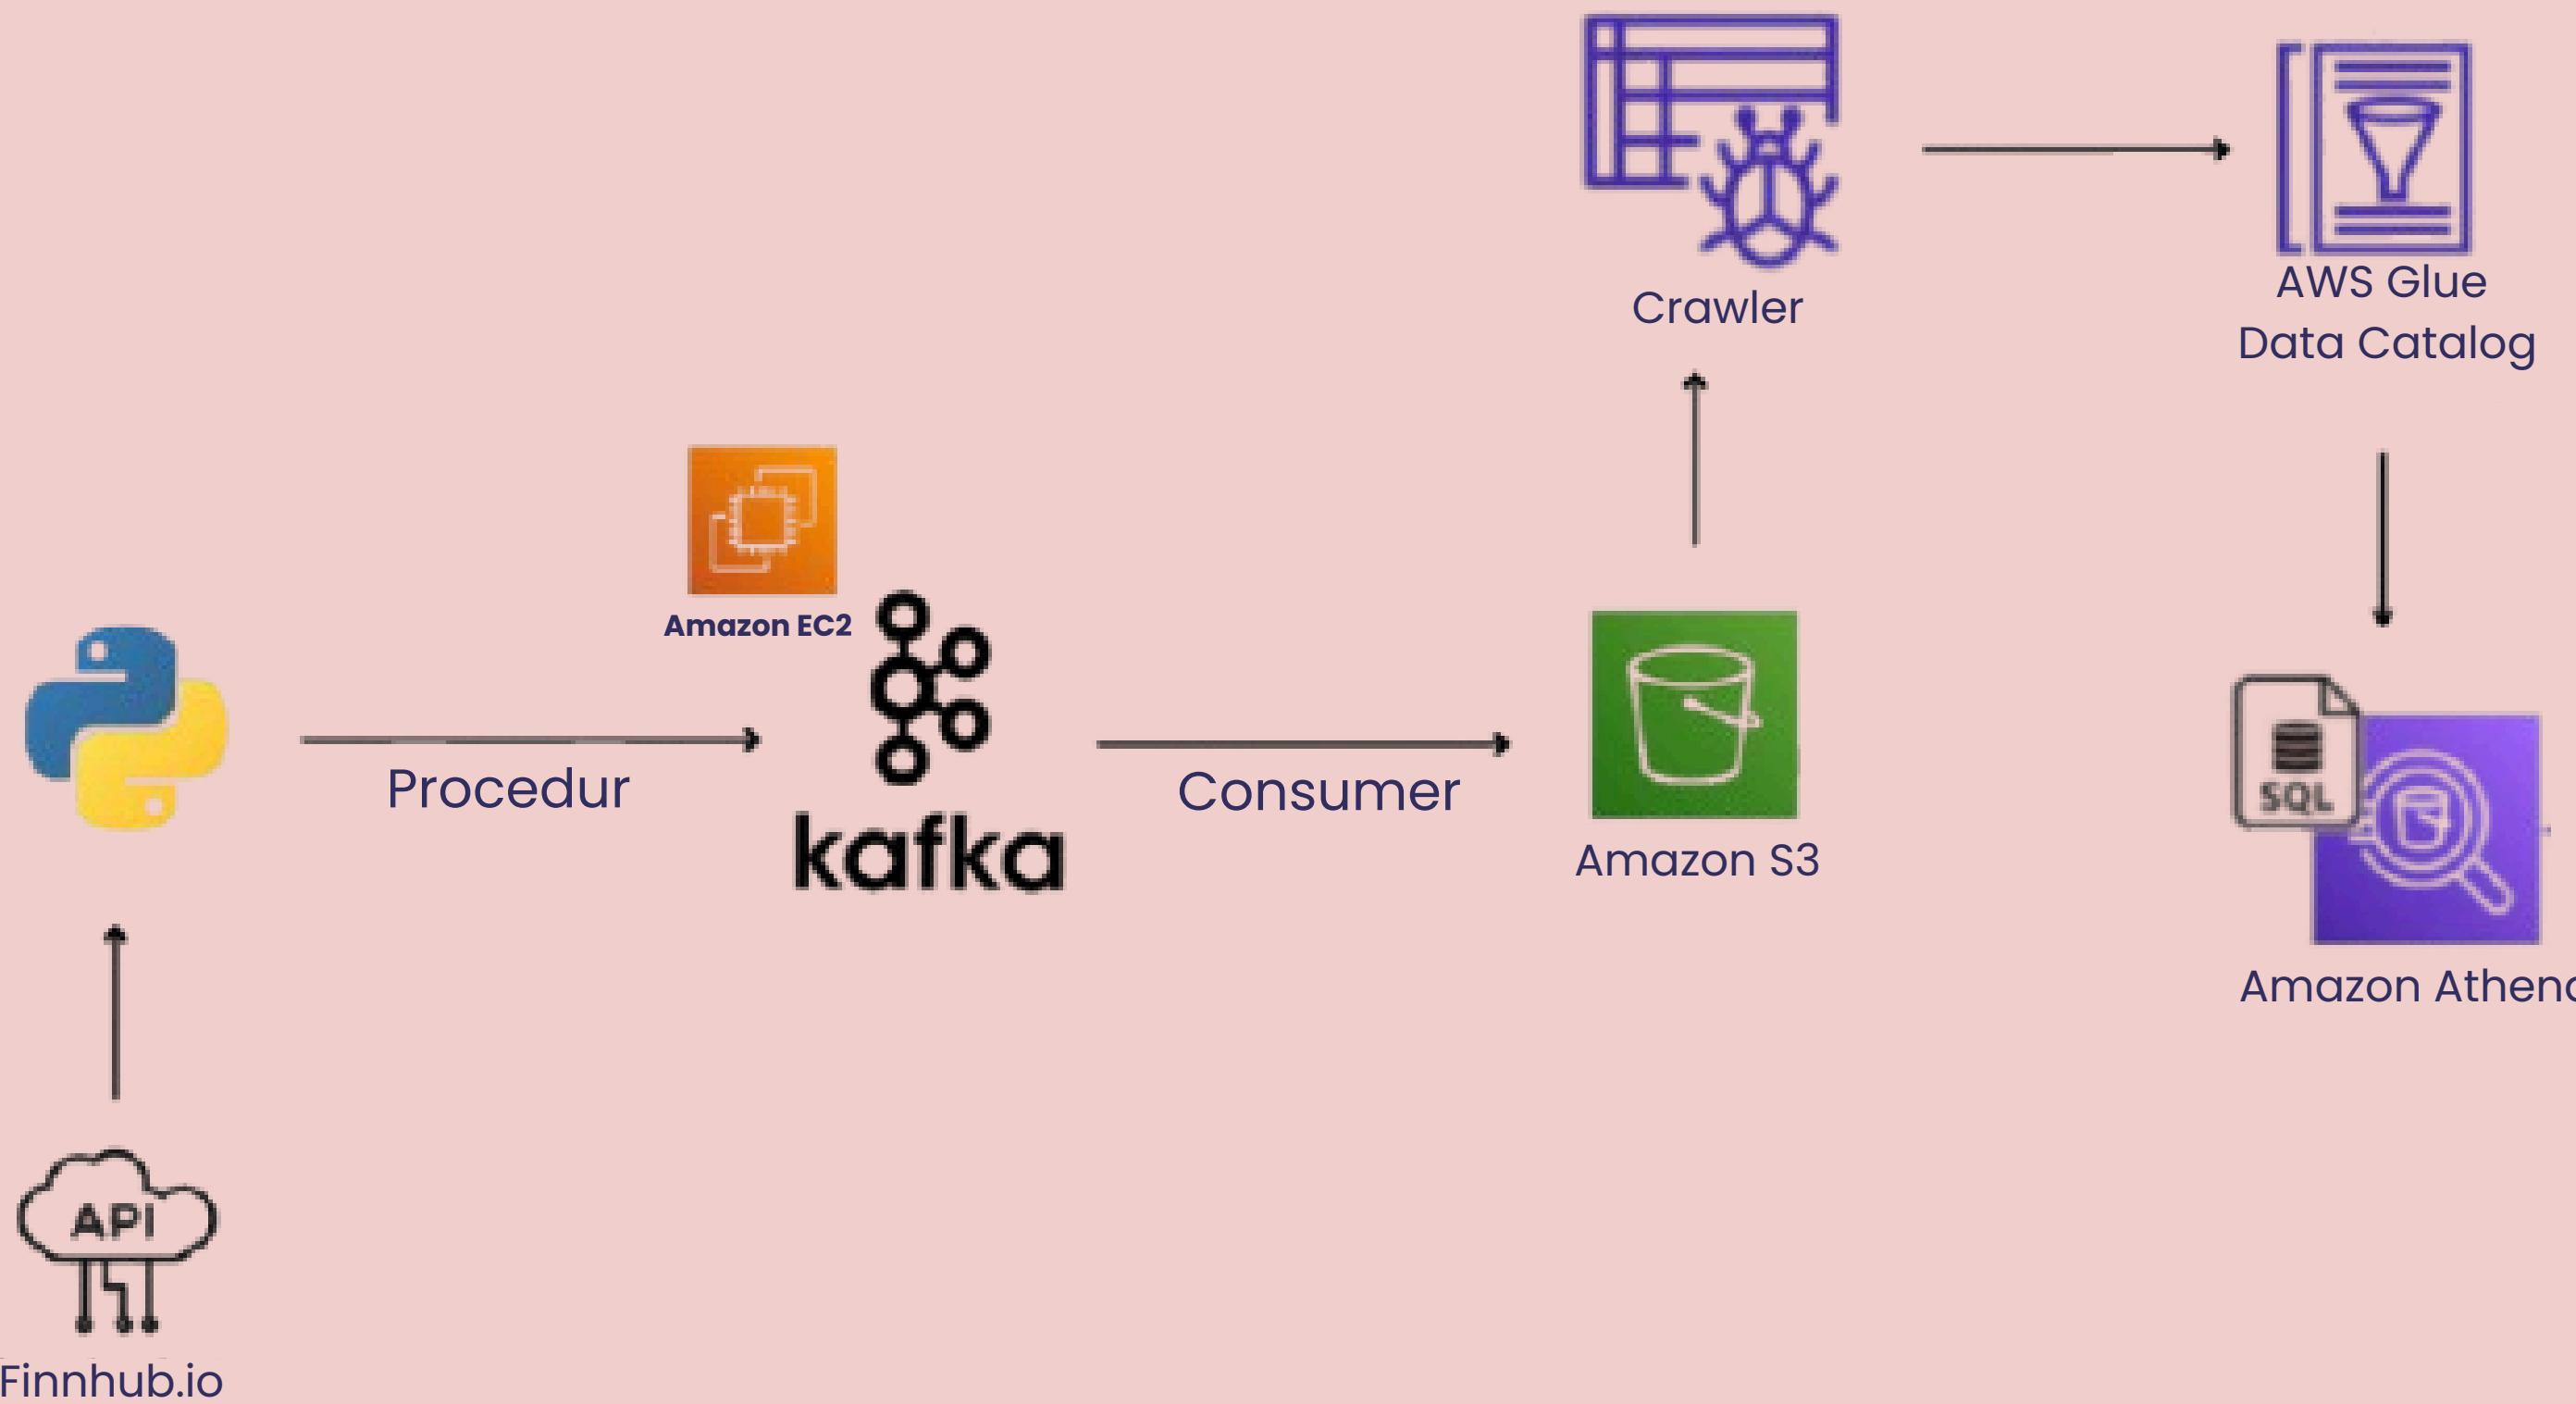
The objective of this project is to build a functional real-time data pipeline for U.S. stock market data.

The system focuses on collecting live stock data from an API, streaming it through Kafka, and storing it in Amazon S3 in a queryable format.

This project aims to demonstrate the end-to-end process of capturing, processing, and preparing real-time data for analysis using Athena.



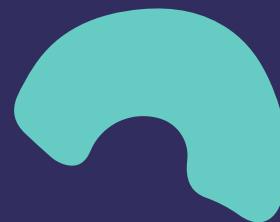
ARCHITECTURE



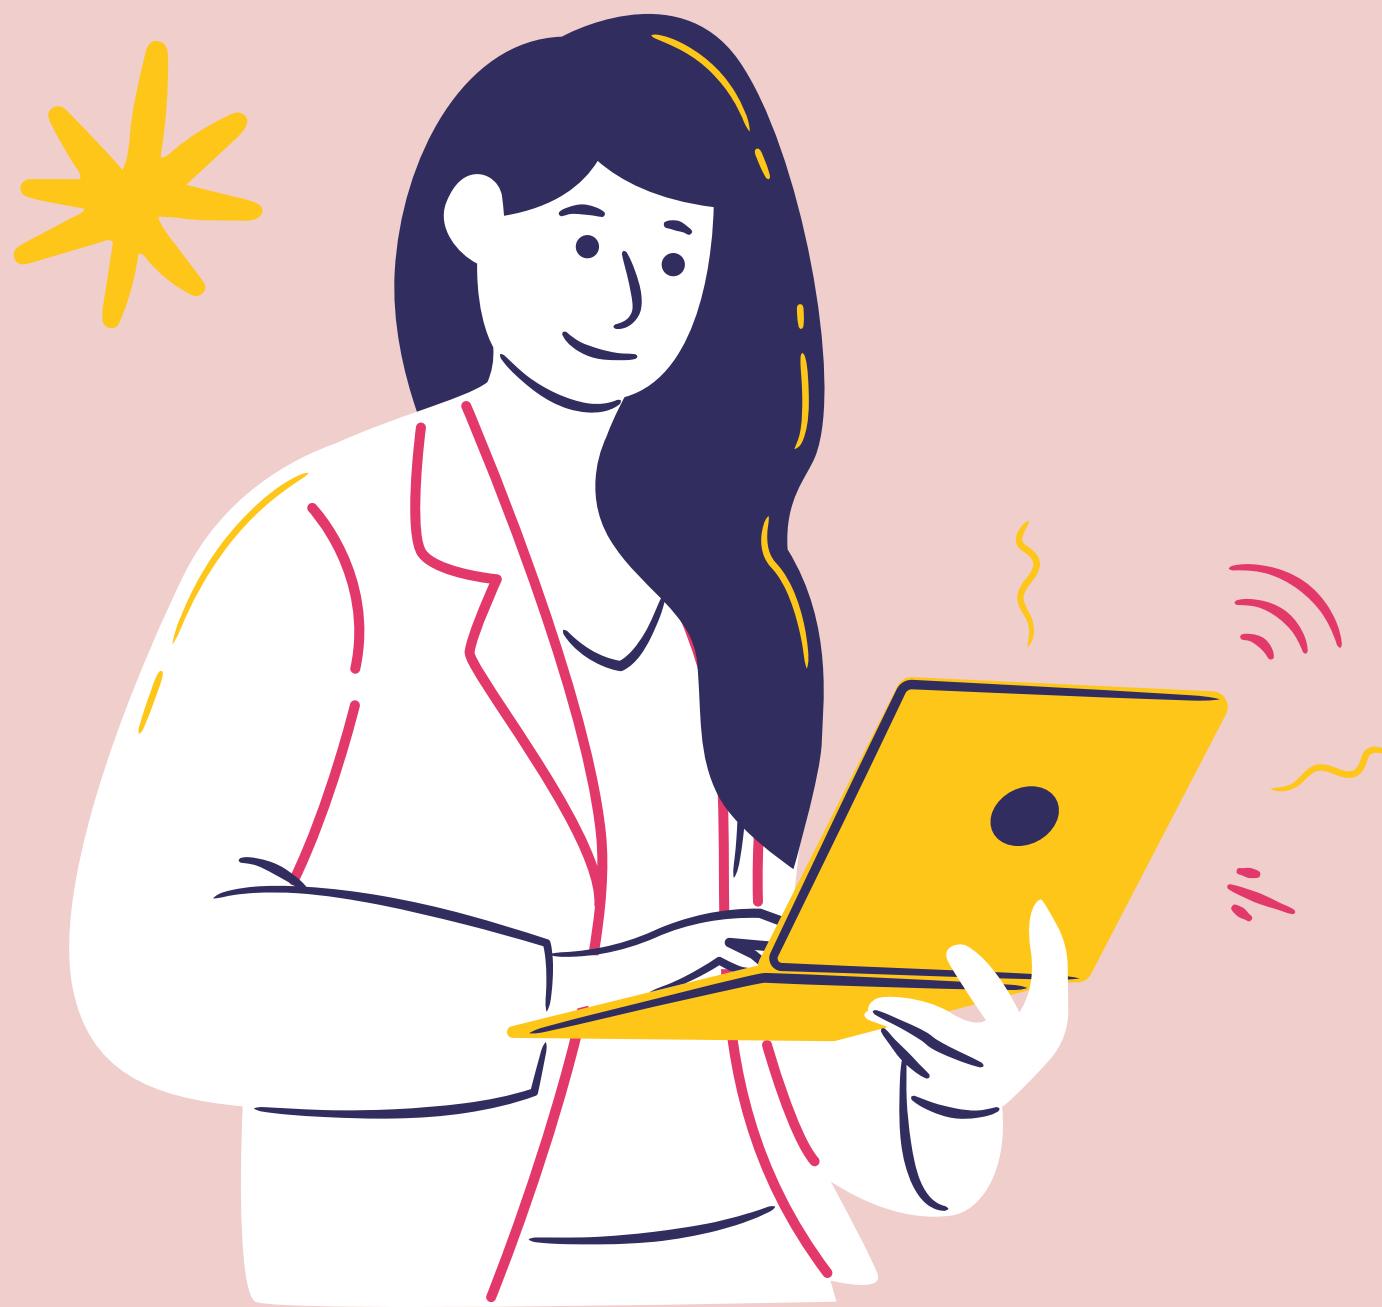
TOOLS & TECHNOLOGIES



- **Kafka** :Used as the message broker to stream live stock market data.
- **Python** : Handles data ingestion from the API and produces messages into Kafka.
- **AWS EC2**: Hosts the Kafka broker and Python producer script.
- **AWS S3** : Stores the streamed data in a scalable, cost-efficient data lake.
- **AWS Glue Crawler** : Automatically detects schema and prepares the data for querying.
- **Amazon Athena**: Runs SQL queries directly on S3 for analysis.



IMPLEMENTATION



REAL-TIME DATA COLLECTION

- Python script connects to Finnhub WebSocket → receives live stock ticks (AAPL, TSLA, MSFT, etc.)

MESSAGE STREAMING WITH KAFKA

- Producer sends each tick into a Kafka topic
- Consumer reads messages and processes them in real time

DATA STORAGE IN AMAZON S3

- Consumer writes validated records as files into S3 buckets
- Folder structure organized by date for easy partitioning

METADATA & SCHEMA WITH AWS GLUE

- Glue Crawler scans the S3 bucket
- Automatically creates / updates tables in AWS Glue Data Catalog

QUERYING WITH AMAZON ATHENA

- Athena reads the Parquet/CSV files using the Glue schema
- Enables SQL analysis directly on S3 data

RESULTS

- Real-time trade data successfully streamed from Finnhub into Kafka.
- Consumer batches the messages and stores them as Parquet files in S3.
- Glue Crawler updates the table schema automatically.
- Athena queries confirm the data is captured and processed correctly.

RESULTS DOCUMENTATION

Procedure run

```
Connected to Finnhub WebSocket
Subscribed: AAPL
Subscribed: GOOGL
Subscribed: TSLA
Subscribed: MSFT
Sent to Kafka: {'c': ['1'], 'p': 415.81, 's': 'TSLA', 't': 1763997375165, 'v': 40}
Sent to Kafka: {'c': ['1', '8'], 'p': 415.89, 's': 'TSLA', 't': 1763997374633, 'v': 40}
Sent to Kafka: {'c': ['1', '8'], 'p': 415.88, 's': 'TSLA', 't': 1763997374633, 'v': 40}
```

Data stored in S3

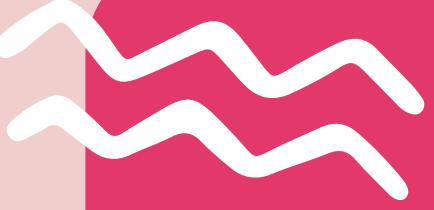
Amazon S3 > Bucket > finnhub-stock-market > stock_data/						
Objek (17)						
	<input type="button" value="C"/>	<input type="button" value="Salin URI S3"/>	<input type="button" value="Salin URL"/>	<input type="button" value="Unduh"/>	<input type="button" value="Buka L"/>	<input type="button" value="Hapus"/>
Objek adalah entitas fundamental yang disimpan di Amazon S3. Anda dapat menggunakan inventaris Amazon S3 untuk mendapatkan daftar semua objek dalam bucket Anda. Agar orang lain dapat mengakses objek Anda, Anda harus memberikan izin secara eksplisit kepada mereka. Pelajari selengkapnya						
#	▼	symbol	▼	minute	▼	avg_price
1		MSFT		2025-11-24 14:48:00.000		470.0875000000003
2		GOOGL		2025-11-24 14:48:00.000		317.2825
3		TSLA		2025-11-24 14:48:00.000		409.0215384615385
4		TSLA		2025-11-24 14:52:00.000		409.3218750000003
5		MSFT		2025-11-24 14:52:00.000		468.5201315789474
6		GOOGL		2025-11-24 14:52:00.000		317.0887500000006
7		AAPL		2025-11-24 14:52:00.000		273.8985714285714
8		GOOGL		2025-11-24 14:53:00.000		317.18467741935484
9		TSLA		2025-11-24 14:53:00.000		410.94158385093175
10		MSFT		2025-11-24 14:53:00.000		469.23449438202243

Consumer run

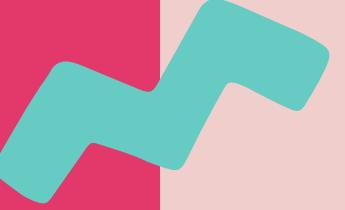
```
Consumer running...
C:\Users\User\AppData\Local\Temp\ipykernel_14744\3217407254.py:23: DeprecationWarning: datetime.d
in a future version. Use timezone-aware objects to represent datetimes in UTC: datetime.datetime.
    timestamp = datetime.datetime.utcnow().strftime("%Y%m%d_%H%M%S")
Uploaded to S3: s3://finnhub-stock-market/stock_data/trades_20251124_151800_dd936e5a_200.parquet
Uploaded to S3: s3://finnhub-stock-market/stock_data/trades_20251124_151901_bb7aa701_200.parquet
```

Query successful in Athena

#	▼	symbol	▼	minute	▼	avg_price	▼	total_volume
1		MSFT		2025-11-24 14:48:00.000		470.0875000000003		1169
2		GOOGL		2025-11-24 14:48:00.000		317.2825		800
3		TSLA		2025-11-24 14:48:00.000		409.0215384615385		1164
4		TSLA		2025-11-24 14:52:00.000		409.3218750000003		1988
5		MSFT		2025-11-24 14:52:00.000		468.5201315789474		3166
6		GOOGL		2025-11-24 14:52:00.000		317.0887500000006		1798
7		AAPL		2025-11-24 14:52:00.000		273.8985714285714		1660
8		GOOGL		2025-11-24 14:53:00.000		317.18467741935484		8894
9		TSLA		2025-11-24 14:53:00.000		410.94158385093175		11566
10		MSFT		2025-11-24 14:53:00.000		469.23449438202243		5737



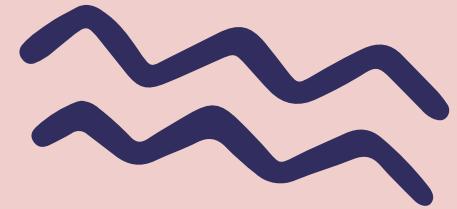
CONCLUSION



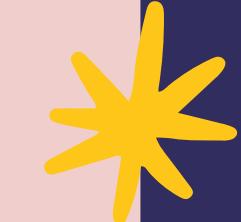
What began as a stream of raw market events moving through Kafka gradually transformed into a well-structured pipeline. Each component played its part: Kafka delivered the data in real time, S3 kept it safely stored, Glue brought structure to the chaos, and

Athena turned it into something we could finally explore and understand. The result is a system that not only handles live market data but makes it meaningful, accessible, and ready for deeper analysis. This pipeline shows how the right tools can turn fast-moving information into clear, actionable insight.





NEXT STEPS / IMPROVEMENTS



ADD MONITORING WITH GRAFANA

Build dashboards to track Kafka throughput, consumer lag, and S3 ingestion for real-time visibility.

OPTIMIZE S3 STRUCTURE WITH PARTITIONING

Organize Parquet files by symbol/date to make Athena queries faster and reduce scan costs.

IMPROVE DATA QUALITY CONTROLS

Add validation and error handling to prevent corrupted or incorrect files from reaching S3.

AUTOMATE METADATA & ETL WITH AWS GLUE

Use crawlers and scheduled jobs to automatically manage schema, partitions, and transformations.



Feel Free to Reach Me Here !

Email : Awinardi1004@gmail.com

LinkedIn : www.linkedin.com/in/winardi-/

GitHub : <https://github.com/awinardi1004>

